

Rohini Dhankhar 2301010266
B.Tech Cse Core-D

Operating System Assignment - 2

PART - A

Ques. 1) Logical Address \rightarrow Physical Address

- CPU generates a logical address (virtual address).
- The Memory Management Unit (MMU) translates it into a physical address.
- Uses page table (in paging) or segment table (in segmentation).

Example: "

Logical Address = Page No., offset

Page Table \rightarrow Frame No.

Physical Address = (Frame No., offset)

Ques. 2) • Internal fragmentation: A 200KB block given to a 180KB process leaves 20KB wasted inside.

- External fragmentation: Free blocks [100KB, 300KB] process requires 250KB \rightarrow cannot fit despite enough total free memory.

• Solution (beyond compaction):

- Paging (removes external fragmentation).

- Segmentation with paging.

- Dynamic allocation (Buddy system, slab allocation).

- Ans 3)
- Memory divided into fixed-size pages (e.g. 4KB).
 - Process is divided into pages & stored in free frames.

• Trade-offs:

- Extra memory overhead (page tables).
- Access speed reduced due to page table lookup, improved with TLB.
- No external fragmentation, but last page may have internal fragmentation.

Ans 4) • Hardware support:

- Page Table Base Register (PTBR)
- Translation Lookaside Buffer (TLB)
- Protection bits in page table entries
- Indirection:
 - If page is present \rightarrow physical address generated
 - If not present \rightarrow page fault \rightarrow OS loads required page from disk \rightarrow updates page table

Ans 5) • Virtual address = 16 bits \rightarrow address space = 2^{16} = 65,536 bytes

- Page size = 1KB = 2^{10} bytes
- No. of virtual pages = $2^{16} / 2^{10} = 2^6 = 64$ pages
- Page table entries = 64, each of 2 bytes \rightarrow Page Table Size = 128 bytes

PART - B

Ans. 6) • Given:

- Total Free memory = 1000 KB

- Processes:

- $P_1 = 212 \text{ KB}$

- $P_2 = 417 \text{ KB}$

- $P_3 = 112 \text{ KB}$

- $P_4 = 426 \text{ KB}$

a) First-Fit Allocation

- $P_1 \rightarrow$ allocated (212 KB) \rightarrow remaining = 788 KB

- $P_2 \rightarrow$ allocated (417 KB) \rightarrow remaining = 371 KB

- $P_3 \rightarrow$ allocated (112 KB) \rightarrow remaining = 259 KB

- $P_4 \rightarrow$ requires 426 KB \rightarrow cannot fit (only 259 KB left)

Unused memory = 259 KB

b) Best-Fit Allocation

- $P_1 \rightarrow$ allocated (212 KB) \rightarrow remaining = 788 KB

- $P_2 \rightarrow$ allocated (417 KB) \rightarrow remaining = 371 KB

- $P_3 \rightarrow$ allocated (112 KB) \rightarrow remaining = 259 KB

- $P_4 \rightarrow$ needs 426 KB \rightarrow cannot fit

Unused memory = 259 KB

c) Worst-Fit Allocation

- $P_1 \rightarrow$ allocated (212 KB) \rightarrow remaining = 788 KB

- $P_2 \rightarrow$ allocated (417 KB) \rightarrow remaining = ³⁷¹259 KB

- $P_3 \rightarrow$ allocated (112 KB) \rightarrow remaining = 259 KB

- $P_4 \rightarrow$ needs 426 KB \rightarrow cannot fit

Unused Memory = 259 KB

Final Result

- First - Fit unused = 259 KB
- Best - Fit unused = 259 KB
- Worst - Fit unused = 259 KB

Conclusion:

All these methods give the same utilization in this case. None can allocated P_4 .

Ex-7 Reference string: 7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 2
Frames: 3

a) FIFO

- Replaces oldest page
- Total 8 page faults

b) Optimal

- Replaces page not needed for longest time
- Total 6 page faults (best)

c) LRU

- Replaces least recently used page
- Total 9 page faults

d) Summary:

- FIFO = 8
- Optimal = 6
- LRU = 9

e) Best performance & Belady's anomaly

- Optimal performs best (minimum faults)
- LRU is a practical approximation
- FIFO can suffer from Belady's anomaly, while Optimal & LRU do not

Ans-8 Given: disk write = 10 ms, memory write = 100 ns,
30% of replaced pages are dirty, 1000 pages
replaced.

- a) Additional time overhead due to dirty pages
- Dirty pages = 30% of 1000 = 300 pages.
 - Time to write one dirty page to disk = 10 ms.
 - Total disk write time = $300 \times 10 \text{ ms} = 3000 \text{ ms} = 3 \text{ s}$.
 - Memory write contribution: $300 \times 100 \text{ ns} = 30,000 \text{ ns} = 0.00003 \text{ s}$ - negligible.

Answer: 3.0 seconds extra overhead (dominant
cost from disk writes).

b) Optimization techniques

- Use write-back with dirty bit (write only when needed).
 - Perform asynchronous / background writes so eviction doesn't block.
 - Batch writes together to reduce I/O overhead.
- These reduce the extra 3 seconds of overhead.

Ans-9) Given: multiple memory-intensive tasks (object
detection - mission-critical, route planning, informa-
-tion - less critical).

- a) Using working set model & page replacement
- Working set: Monitor each task's working set (recent pages it actively uses).
 - Guarantee frames to critical task: Reserve enough frames for object detection equal to its working set so it won't thrash.
 - Replacement policy: Use Working-Set-aware LRU (or true Working Set algo) that keeps pages for tasks whose working set is active; prefer

- existing pages of low priority tasks (infotainment)
- Result: Object detection keeps required pages in memory (low page faults); less critical tasks adapt or are throttled.

b) Memory allocation strategy

- Priority-based dynamic allocation with quotas
- Reserve a minimum guaranteed frame quota for real-time, safety tasks (object detection, sensor fusion)
- Allow remaining frames to be shared dynamically using a global LRU / Working Set policy
- If memory pressure rises, demote or throttle non-critical tasks (reduce their quota / lower fidelity)
- Justification: Ensures real-time responsiveness & avoid thrashing for critical tasks while keeping overall memory utilization efficient & fair