Operating System Assignment - 4

Rohini
2301010266

1. Distributed Deadlocks Detection Simulation -

→ Fragments :

- S_1 : $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4$
- S_2 : $P_2 \rightarrow P_5$, $P_5 \rightarrow P_6$
- S_3 : $P_6 \rightarrow P_1$

a) Global wait - for Graph (combined)

$P_1 \rightarrow P_2 \rightarrow P_5 \rightarrow P_6 \rightarrow P_1$ (cycle) Also $P_3 \rightarrow P_4$ (separate)

b) Deadlock ?

→ Yes Process Involved in dead lock cycle

P_1, P_2, P_5, P_6

c) Suggested distributed algorithm

Use the chandy - Misra - Mass edge - chasing (Probe) algorithm for distributed deadlock detection each site send Probes along wait - for edges to detect cycles without centralized Graph assembly.

2. Distributed File System Performance :

Given : local = 5ms, remote = 25 ms

Prob (remote) = 0.3

a) Excepted access time :

$E = 0.7 \times 5ms + 0.3 \times 25 ms = 3.5 + 37.5 = 11ms.$

b) The caching Strategy :

Client - Side Read cache with LRU + TTL - based validation

→ Justification: Frequently Read Remark files will be served locally reducing Remote access (0.3 fraction) LRU evicts less-used items TTL keeps staleness bounded. Implanes average intency while keeping consistency manageable.

3. Checkpointing Mix to Meet RPO ≤ 1s

Sol Given,

Full = Rooms, incremental = 50 ms, RPO = 1s

a) Proposed mix (over 10s)

- Take one full checkpoint every 10s (at t = 0 in the period)
- Take incremental checkpoint every 2s (at t = 4, 2)

Total overhead (Per loss) : 1 × 200 ms + 9 × 50 ms
= 200 + 450 = 650 ms.

b) Reasoning.

- With incremental every 1s, the maximum work lost on failure < 1s → meets → RPO
- Full once per 10s bound Recovery time 1s restricting an older full + a small no of incrementals is feasible)
- This mix minimize full checkpoint cost while keeping incrementals frequently high enough to meet RPO.

4. Case Study - Global E-commerce Platform.

a) Distributed scheduling challenges for flash sales.

- Massive sudden spikes in Request.

- Geographic distribution & latency data locality matters for latency & inventory correctness
- Heterogeneous nodes

b) Fault - tolerance Strategy (RTO & RPO):
- Active - active multiple - region development:
Service Run concurrently in multiple Regions so failure is seamless.
(RTO ≈ near - zero at service load)

- Data Strategy:
→ Critical transactional data: Use synchronous replication within the region to guarantee consistency & low RPO; cross Region replication but with frequent replication to keep RPO small
→ Data lag / less - critical data: Uses multi - region eventual consistency with frequent async replication & local caches.

- Recovery Plans: Automated failure by health check & Global load balance, warm stand by replicas ready to accept traffic, & durable messaging to buffer incoming requests during failure
- Operational measures: chaos testing, backups, runbook & automated scaling to minimize RTO, use snapshots & incremental backups to bound RPO

Sunil 21/11/25